

Question:

1. **Dataset Selection:** Download any publicly available dataset suitable for **multiclass classification**. Examples include the Iris dataset, Wine dataset, or any other dataset with **more than two classes**. Make sure the dataset has enough samples and features for training a logistic regression model.

2. **Data Preparation:**

- Load the dataset and perform any necessary preprocessing
- Split the dataset into training and testing sets (e.g., 70% training, 30% testing).

3. **Model Building:**


- Build a **multiclass logistic regression** model using Python (e.g., scikit-learn).
- Train the model on the training set.

4. **Evaluation:**

- Evaluate the model on the testing set using appropriate **performance metrics** for multiclass classification, including:
 - Accuracy
 - Precision (per class and macro/micro average)
 - Recall (per class and macro/micro average)
 - F1-score (per class and macro/micro average)
 - Confusion matrix

5. **Analysis:**

- Display all the above metric values clearly.
- Provide a brief interpretation of the model's performance. For example, mention which classes are predicted well and which classes are misclassified more often.



```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('/Users/ishan-college/Desktop/AI (ML Course)/Datasets/animes.csv')

df = df.dropna(subset=["ranked"]) # ranked must not be NaN
df["ranked"] = pd.to_numeric(df["ranked"], errors="coerce")
df = df.dropna(subset=["ranked"])

def convert_rank(r):
    if r <= 200:
        return "Top"
    elif r <= 1000:
        return "Mid"
    else:
        return "Low"

df["rank_category"] = df["ranked"].apply(convert_rank)
target_col = "rank_category"

features = ["members", "popularity", "episodes", "score"]
df = df.dropna(subset=features)

X = df[features]
y = df[target_col]

le = LabelEncoder()
y = le.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LogisticRegression(
    max_iter=5000,
    solver='lbfgs',
    multi_class='auto',
    class_weight='balanced'
)
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=le.classes_, yticklabels=le.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

print("\nConfusion Matrix:")
print(cm)

print("\nClassification Report:")
print(classification_report(y_test, y_pred, target_names=le.classes_))
|
example_idx = 0
example_features = X_test_scaled[example_idx].reshape(1, -1)
predicted_class = le.inverse_transform(model.predict(example_features))[0]
print("\nExample Prediction:")
print("Predicted class =", predicted_class)

```